

Capital Bike Share

March 24, 2019

0.1 Introduction

Bike sharing systems are a means of renting bicycles where the process of obtaining membership, rental, and bike return is automated via a network of kiosk locations throughout a city. Using these systems, people are able rent a bike from a one location and return it to a different place on an as-needed basis. Currently, there are over 500 bike-sharing programs around the world.

The data generated by these systems makes them attractive for researchers because the duration of travel, departure location, arrival location, and time elapsed is explicitly recorded. Bike sharing systems therefore function as a sensor network, which can be used for studying mobility in a city. In this competition, participants are asked to combine historical usage patterns with weather data in order to forecast bike rental demand in the Capital Bikeshare program in Washington, D.C.

Evaluation: Submissions are evaluated one the Root Mean Squared Logarithmic Error (RMSLE

0.1.1 Data

datetime - hourly date + timestamp

season - 1 = spring, 2 = summer, 3 = fall, 4 = winter

holiday - whether the day is considered a holiday

workingday - whether the day is neither a weekend nor holiday

weather - 1: Clear, Few clouds, Partly cloudy, Partly cloudy

2: Mist + Cloudy, Mist + Broken clouds, Mist + Few clouds, Mist

3: Light Snow, Light Rain + Thunderstorm + Scattered clouds, Light Rain + Scattered clouds

4: Heavy Rain + Ice Pellets + Thunderstorm + Mist, Snow + Fog

temp - temperature in Celsius

atemp - "feels like" temperature in Celsius

humidity - relative humidity

windspeed - wind speed

casual - number of non-registered user rentals initiated

registered - number of registered user rentals initiated

count - number of total rentals

Importing libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from pandas.io.json import json_normalize
```

```

import scipy.stats as stats
import pylab
from numpy.random import seed
from numpy.random import randn
from scipy.stats import shapiro
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_log_error
%matplotlib inline

```

Importing datasets

```

In [2]: train = pd.read_csv("train.csv")
        test = pd.read_csv("test.csv")

```

```

In [3]: print(f'Shape of train dataset: {str(train.shape)}')
        print(f'Shape of test dataset: {str(test.shape)}')

```

Shape of train dataset: (10886, 12)

Shape of test dataset: (6493, 9)

```

In [4]: train.head()

```

```

Out[4]:
      datetime  season  holiday  workingday  weather  temp  atemp  \
0  2011-01-01 00:00:00      1      0          0        1   9.84  14.395
1  2011-01-01 01:00:00      1      0          0        1   9.02  13.635
2  2011-01-01 02:00:00      1      0          0        1   9.02  13.635
3  2011-01-01 03:00:00      1      0          0        1   9.84  14.395
4  2011-01-01 04:00:00      1      0          0        1   9.84  14.395

      humidity  windspeed  casual  registered  count
0           81          0.0        3          13      16
1           80          0.0        8          32      40
2           80          0.0        5          27      32
3           75          0.0        3          10      13
4           75          0.0        0           1       1

```

```

In [5]: train.describe()

```

```

Out[5]:
      season  holiday  workingday  weather  temp  \
count  10886.000000  10886.000000  10886.000000  10886.000000  10886.000000
mean      2.506614      0.028569      0.680875      1.418427      20.23086
std      1.116174      0.166599      0.466159      0.633839      7.79159
min      1.000000      0.000000      0.000000      1.000000      0.82000
25%      2.000000      0.000000      0.000000      1.000000      13.94000
50%      3.000000      0.000000      1.000000      1.000000      20.50000
75%      4.000000      0.000000      1.000000      2.000000      26.24000
max      4.000000      1.000000      1.000000      4.000000      41.00000

```

	atemp	humidity	windspeed	casual	registered \
count	10886.000000	10886.000000	10886.000000	10886.000000	10886.000000
mean	23.655084	61.886460	12.799395	36.021955	155.552177
std	8.474601	19.245033	8.164537	49.960477	151.039033
min	0.760000	0.000000	0.000000	0.000000	0.000000
25%	16.665000	47.000000	7.001500	4.000000	36.000000
50%	24.240000	62.000000	12.998000	17.000000	118.000000
75%	31.060000	77.000000	16.997900	49.000000	222.000000
max	45.455000	100.000000	56.996900	367.000000	886.000000

	count
count	10886.000000
mean	191.574132
std	181.144454
min	1.000000
25%	42.000000
50%	145.000000
75%	284.000000
max	977.000000

```
In [6]: def generateColumnInfo(df):
        cls = []
        nullCount = []
        nonNullCount = []
        nullsPct = []
        uniqCount = []
        dataType = []
        for i,col in enumerate(df.columns):
            cls.append(col)
            nullCount.append(df[col].isnull().sum())
            nonNullCount.append(len(df)-df[col].isnull().sum())
            nullsPct.append((df[col].isnull().sum()*(100)/len(df))
            uniqCount.append(df[col].nunique())
            dataType.append(df[col].dtype)

        column_info = pd.DataFrame(
            {'ColumnName': cls,
             'NullCount': nullCount,
             'NonNullCount': nonNullCount,
             'NullPercent': nullsPct,
             'UniqueValueCount': uniqCount,
             'DataType': dataType
            })
        return(column_info)
```

```
In [7]: generateColumnInfo(train)
```

Out [7]:	ColumnName	NullCount	NonNullCount	NullPercent	UniqueValueCount \
0	datetime	0	10886	0.0	10886

1	season	0	10886	0.0	4
2	holiday	0	10886	0.0	2
3	workingday	0	10886	0.0	2
4	weather	0	10886	0.0	4
5	temp	0	10886	0.0	49
6	atemp	0	10886	0.0	60
7	humidity	0	10886	0.0	89
8	windspeed	0	10886	0.0	28
9	casual	0	10886	0.0	309
10	registered	0	10886	0.0	731
11	count	0	10886	0.0	822

	DataType
0	object
1	int64
2	int64
3	int64
4	int64
5	float64
6	float64
7	int64
8	float64
9	int64
10	int64
11	int64

In [8]: generateColumnInfo(test)

Out [8]:	ColumnName	NullCount	NonNullCount	NullPercent	UniqueValueCount	DataType
0	datetime	0	6493	0.0	6493	object
1	season	0	6493	0.0	4	int64
2	holiday	0	6493	0.0	2	int64
3	workingday	0	6493	0.0	2	int64
4	weather	0	6493	0.0	4	int64
5	temp	0	6493	0.0	49	float64
6	atemp	0	6493	0.0	65	float64
7	humidity	0	6493	0.0	79	int64
8	windspeed	0	6493	0.0	27	float64

Train contains 3 additional columns : casual, registered and count. There are no null values. We will start by converting the text column "datetime" to datetime type. We will extract the date related columns and drop the original column from the train set. We need to retain the datetime column on the test set as it is required to submit the results.

```
In [9]: train['date'] = pd.to_datetime(train['datetime'],format='%Y-%m-%d')
        train.drop(['datetime'],axis=1,inplace=True)
        train['date_year'],train['date_month'],train['date_day'],train['date_weekday'],train['date_weekday_name'] = train['date'].dt.weekday_name
```

```
test['date'] = pd.to_datetime(test['datetime'],format='%Y-%m-%d')
test['date_year'],test['date_month'],test['date_day'],test['date_weekday'],test['date_
```

```
In [10]: train.head()
```

```
Out[10]:
```

	season	holiday	workingday	weather	temp	atemp	humidity	windspeed	\
0	1	0	0	1	9.84	14.395	81	0.0	
1	1	0	0	1	9.02	13.635	80	0.0	
2	1	0	0	1	9.02	13.635	80	0.0	
3	1	0	0	1	9.84	14.395	75	0.0	
4	1	0	0	1	9.84	14.395	75	0.0	

	casual	registered	count	date	date_year	date_month	\
0	3	13	16	2011-01-01 00:00:00	2011	1	
1	8	32	40	2011-01-01 01:00:00	2011	1	
2	5	27	32	2011-01-01 02:00:00	2011	1	
3	3	10	13	2011-01-01 03:00:00	2011	1	
4	0	1	1	2011-01-01 04:00:00	2011	1	


	date_day	date_weekday	date_hour	weekday_name
0	1	5	0	Saturday
1	1	5	1	Saturday
2	1	5	2	Saturday
3	1	5	3	Saturday
4	1	5	4	Saturday

0.1.2 Distribution of dependant variable

We see many outliers (3CEIQR) in the boxplot.


```
In [11]: sns.boxplot(train['count'],orient='v')
```

```
Out[11]: <matplotlib.axes._subplots.AxesSubplot at 0xfb0ed9470>
```




```
In [12]: sns.distplot(train['count'])
```

```
Out[12]: <matplotlib.axes._subplots.AxesSubplot at 0xfb1219c18>
```



```
In [13]: measurements = np.random.normal(loc = 20, scale = 5, size=100)
stats.probplot(train['count'], dist="norm", plot=pylab)
pylab.show()
```



The 'count' variable has a right skew. Quantile-Quantile plot shows that it is not showing normal distribution. Shapiro-Wilk's test confirms this behavior.

```
In [14]: #ShapiroWilk test
stat, p = shapiro(train['count'])
print('p=%.3f' % (p)) #Normal if p>0.1
```


p=0.000

D:\Anaconda\lib\site-packages\scipy\stats\morestats.py:1310: UserWarning: p-value may not be accurate for N > 5000.
warnings.warn("p-value may not be accurate for N > 5000.")

Log transformation of the dependant variable improves the distribution. I will use the log transformation for statistical modeling.

```
In [15]: sns.distplot(np.log1p(train['count']))
```

```
Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0xf155d780>
```




0.2 Exploratory data analysis (EDA)

```
In [16]: #creating additional dataframe for EDA
byUserType = train.drop(['count'],axis=1)
byUserType = pd.melt(byUserType,id_vars = ['season', 'holiday', 'workingday', 'weather',
      'humidity', 'windspeed', 'date',
      'date_year', 'date_month', 'date_day', 'date_weekday', 'date_hour',
      'weekday_name'],value_vars = ['casual', 'registered'],var_name='user_type',val
```


I will analyze each column in this section. I will start with the date/time based columns.

0.2.1 Date Columns

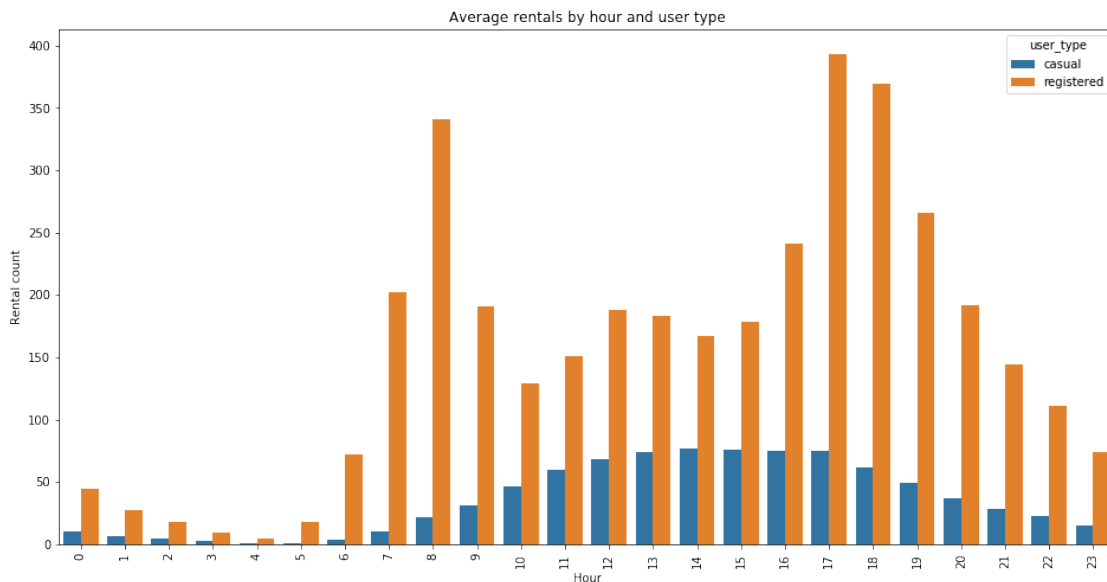
```
In [17]: plt.figure(figsize=(12,7))
sns.boxplot(data=train, x='date_hour',y='count')
plt.title('Average rentals per hour')
plt.xlabel('Hour')
plt.ylabel('Rental count')
plt.show()
```

```
In [18]: plt.figure(figsize=(12,7))
sns.barplot(data=train,x='date_hour',y='count',ci=None)
plt.title('Average rentals per hour')
plt.xlabel('Hour')
plt.ylabel('Rental count')
plt.show()
```



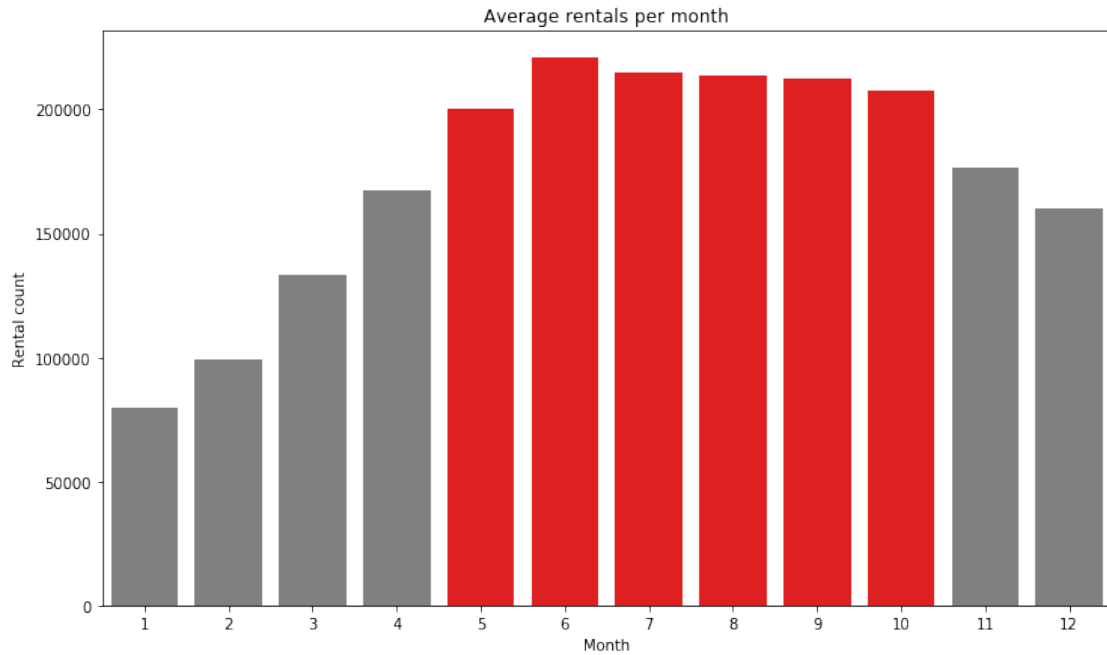
```
In [19]: plt.figure(figsize=(16,8))
plt.xticks(rotation=90)
ax = sns.barplot(data=byUserType,x='date_hour',y='user_count',hue='user_type',ci=None)
plt.title('Average rentals by hour and user type')
plt.xlabel('Hour')
plt.ylabel('Rental count')
plt.show()
```



Demand is relatively high between 7am-9am and 4pm-7pm. This high demand is fueled by registered users. The demand from casual users is high between 11am-5pm. This could be attributed to the visitors and it coincides with the timings of major tourist attractions in the DC area.

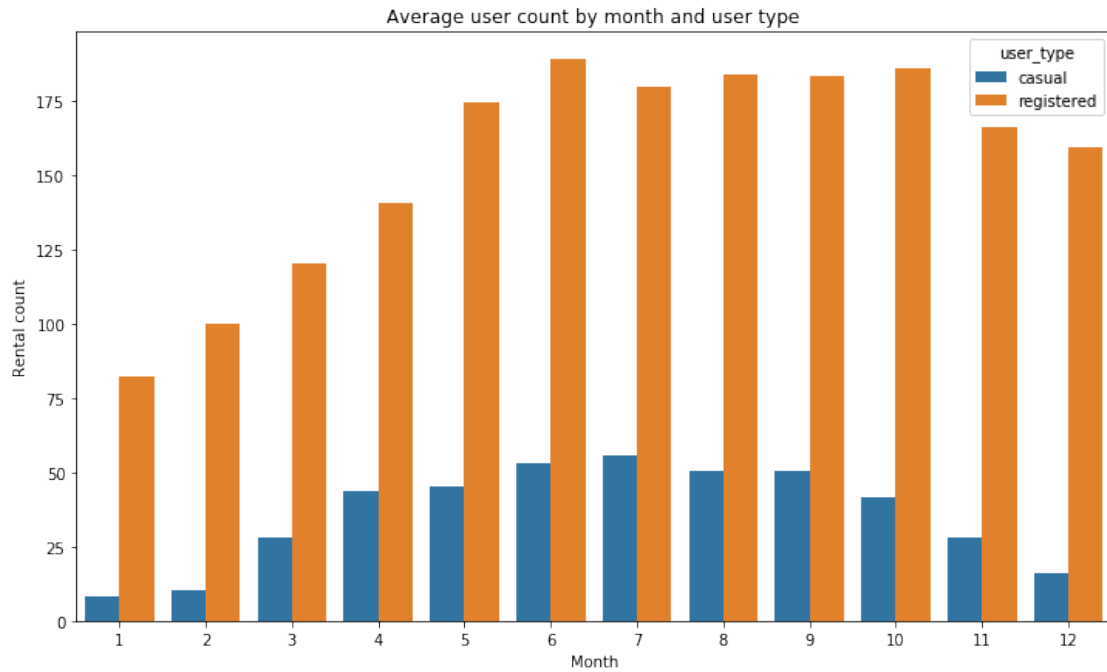
```
In [20]: plt.figure(figsize=(12,7))
countByMonth = train.groupby(['date_month'],as_index=False).sum()[['date_month','count']]
ids = countByMonth['date_month']
values = countByMonth['count']
plt.figure(figsize=(12,7))
clrs = ['grey' if (x < 180000) else 'red' for x in values]
sns.barplot(x=ids,y=values,ci=None,palette=clrs)
plt.title('Average rentals per month')
plt.xlabel('Month')
plt.ylabel('Rental count')
plt.show()
```

<Figure size 864x504 with 0 Axes>

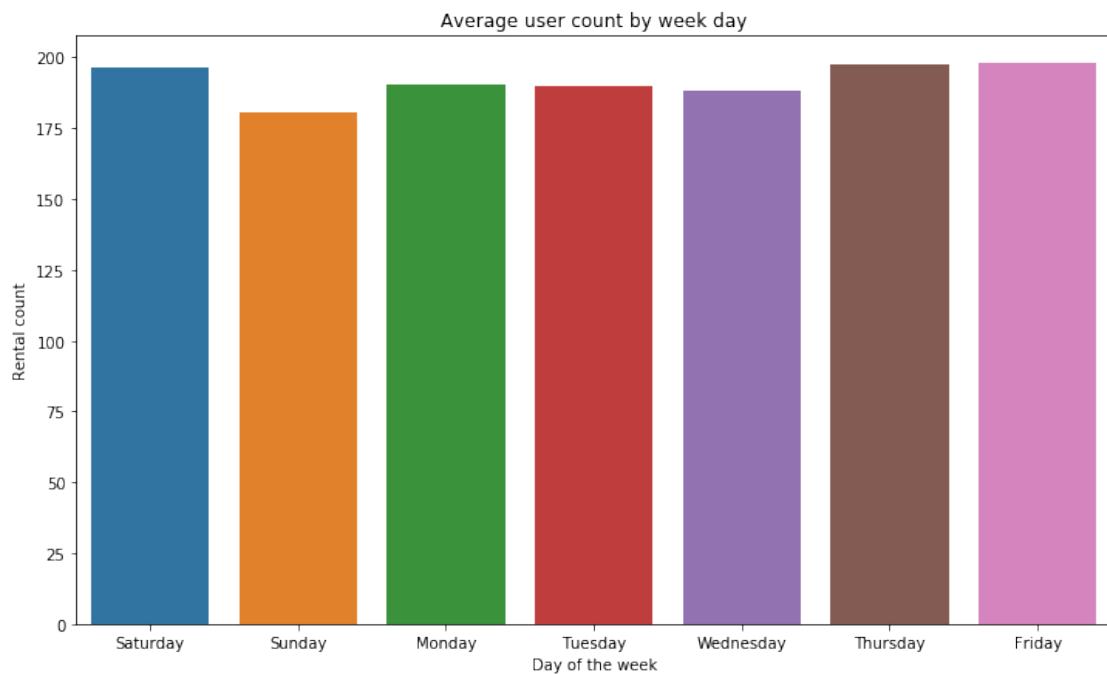


Demand is high during the summer months (May to October). Both registered and casual riders show the same behavior.

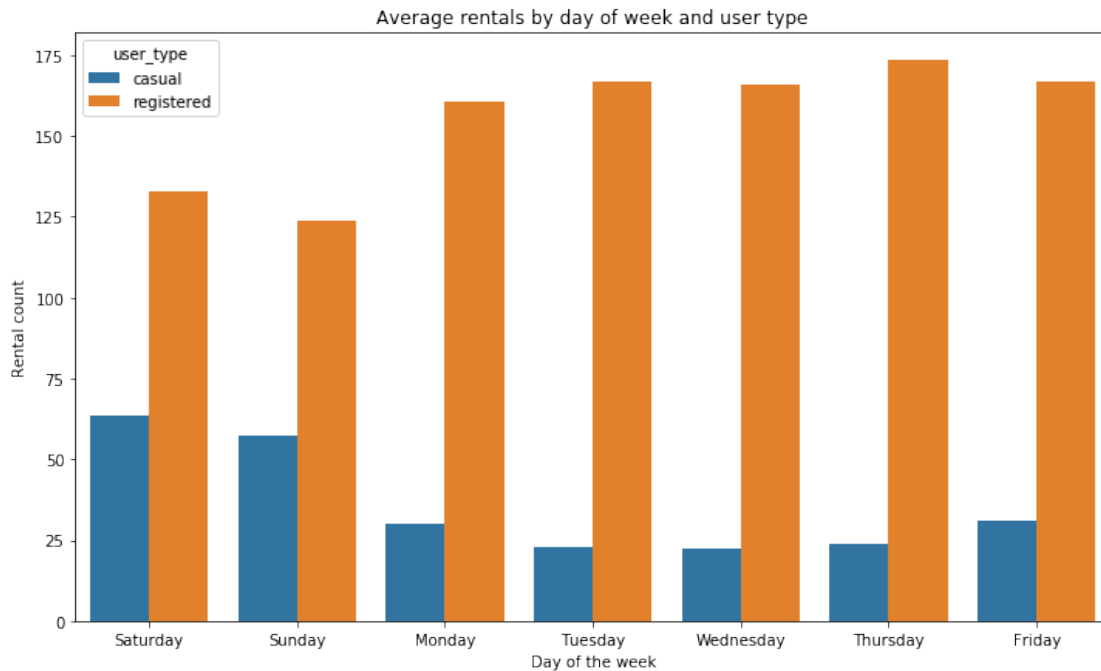
```
In [21]: plt.figure(figsize=(12,7))
sns.barplot(data=byUserType,x='date_month',y='user_count',hue='user_type',ci=None)
plt.title('Average user count by month and user type')
plt.xlabel('Month')
plt.ylabel('Rental count')
plt.show()
```



```
In [22]: plt.figure(figsize=(12,7))
sns.barplot(data=train,x='weekday_name',y='count',ci=None)
plt.title('Average user count by week day')
plt.xlabel('Day of the week')
plt.ylabel('Rental count')
plt.show()
```

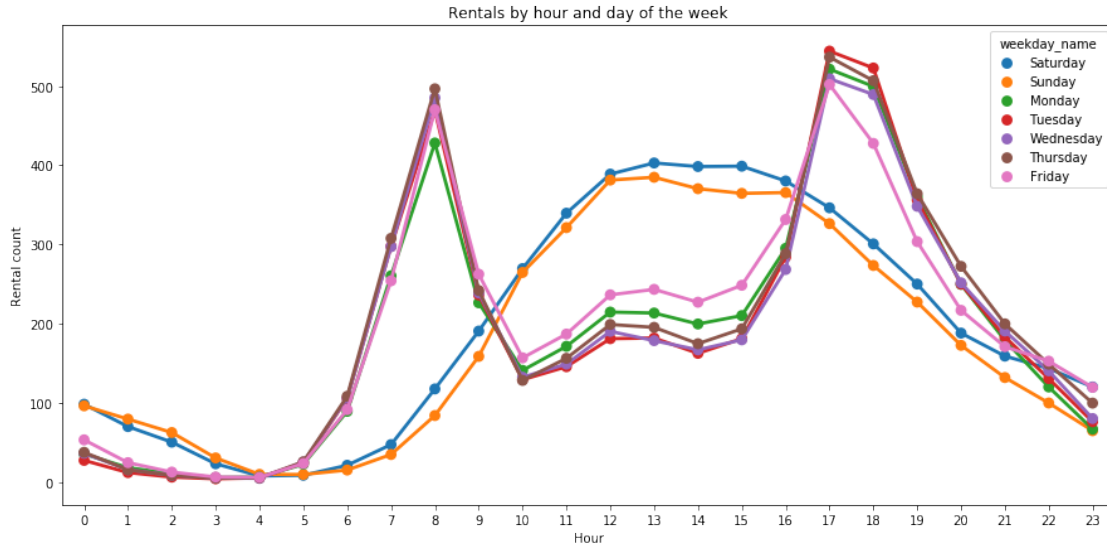


```
In [23]: plt.figure(figsize=(12,7))
sns.barplot(data=byUserType,x='weekday_name',y='user_count',hue='user_type',ci=None)
plt.title('Average rentals by day of week and user type')
plt.xlabel('Day of the week')
plt.ylabel('Rental count')
plt.show()
```



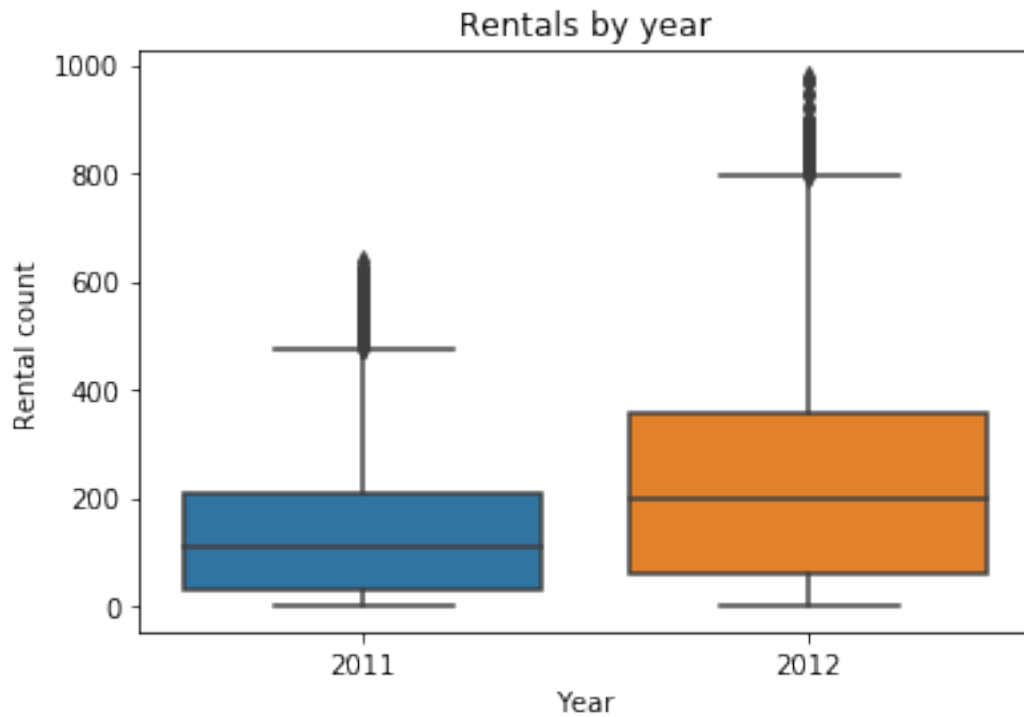
- Demand from registered users is higher on weekdays.
- Demand from casual users is higher on weekends.

```
In [24]: plt.figure(figsize=(15,7))
sns.pointplot(data=train,x='date_hour',y='count',hue='weekday_name',ci=None)
plt.title('Rentals by hour and day of the week')
plt.xlabel('Hour')
plt.ylabel('Rental count')
plt.show()
```



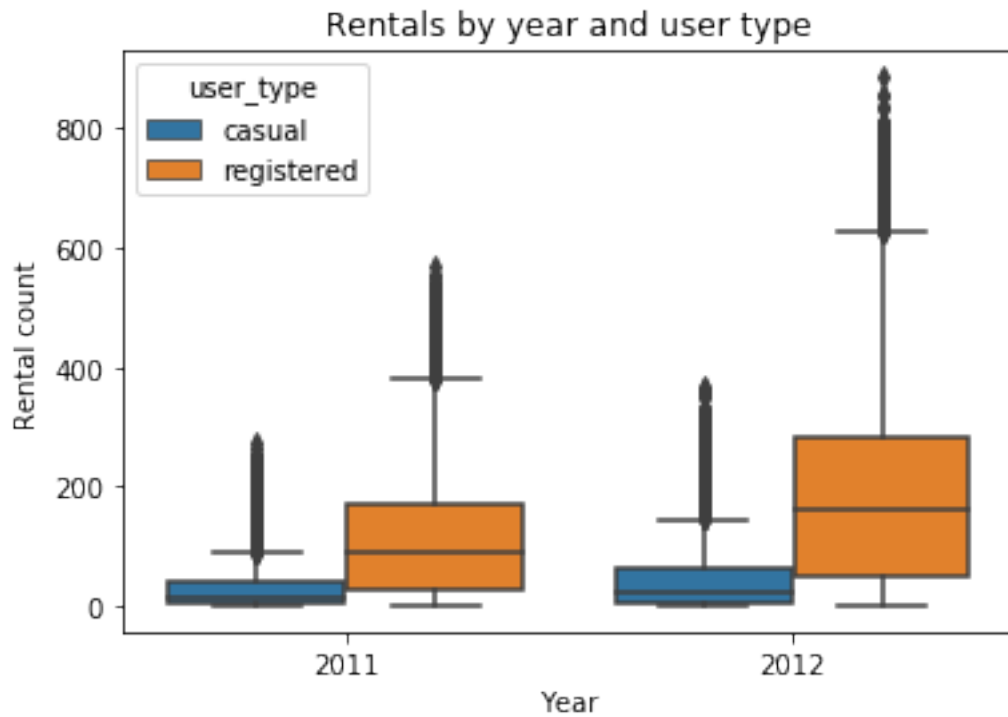
- There is a significant difference in the demand between weekdays and weekends.
- On weekdays, the peak demand is during commuting hours (7am-9am) and (4pm-7pm).
- On weekends, the peak demand is between noon and 4pm.

```
In [25]: sns.boxplot(data=train,x='date_year',y='count')
plt.title('Rentals by year')
plt.xlabel('Year')
plt.ylabel('Rental count')
plt.show()
```



```
In [26]: sns.boxplot(data=byUserType,x='date_year',y='user_count',hue="user_type")
plt.title('Rentals by year and user type')
plt.xlabel('Year')
plt.ylabel('Rental count')
```

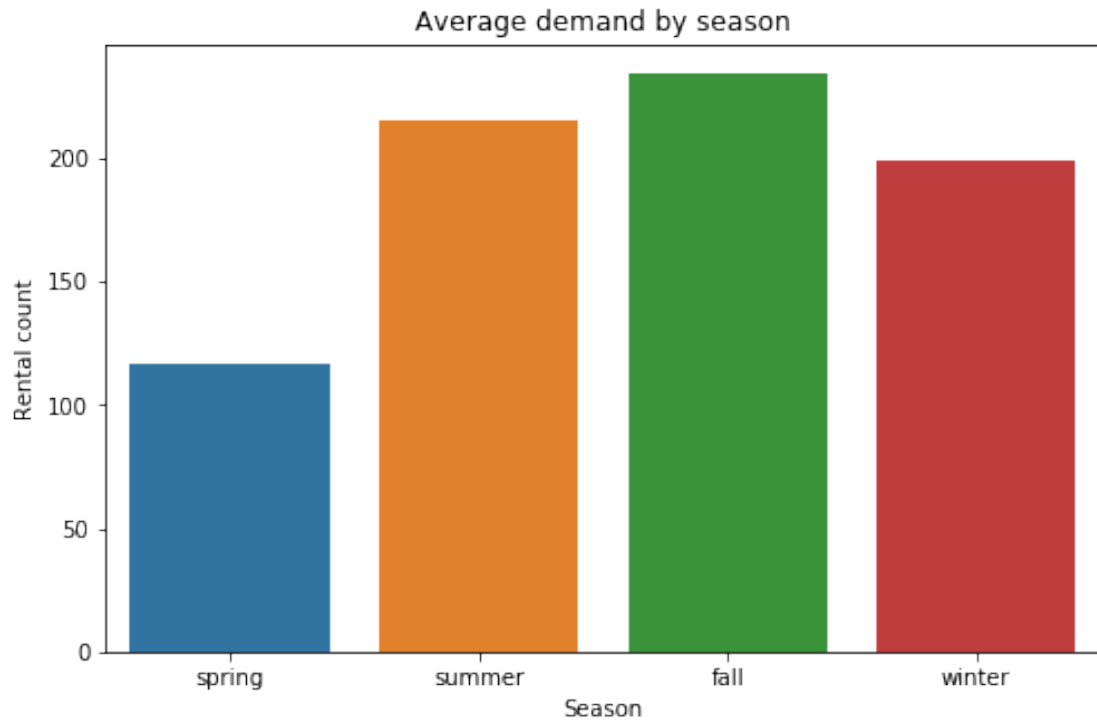
```
Out[26]: Text(0,0.5,'Rental count')
```



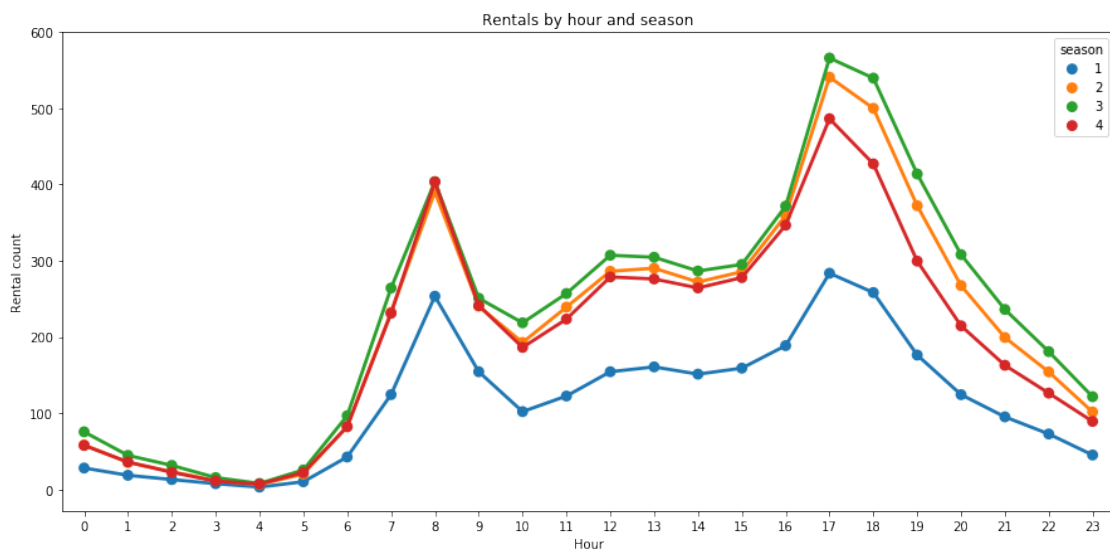
0.2.2 Season

```
In [27]: plt.figure(figsize=(8,5))
sns.barplot(data=train,x='season',y='count',ci=None)
ax = plt.axes()
plt.xticks(rotation=0)
ax.set_xticklabels(['spring','summer','fall','winter'])
plt.title('Average demand by season')
plt.xlabel('Season')
plt.ylabel('Rental count')
plt.show()
```

D:\Anaconda\lib\site-packages\matplotlib\cbook\deprecation.py:107: MatplotlibDeprecationWarning
warnings.warn(message, mplDeprecation, stacklevel=1)




```
In [28]: plt.figure(figsize=(15,7))
g = sns.pointplot(data=train,x='date_hour',y='count',hue='season',ci=None)
plt.title('Rentals by hour and season')
plt.xlabel('Hour')
plt.ylabel('Rental count')
plt.show()
```




0.2.3 Holiday and Working Days

```
In [29]: plt.figure(figsize=(8,5))
sns.boxplot(data=byUserType, x='holiday', y='user_count', hue='user_type')
plt.xlabel('Holiday')
plt.ylabel('Rental count')
plt.show()
```



```
In [30]: plt.figure(figsize=(8,5))
sns.boxplot(data=byUserType, x='workingday', y='user_count', hue='user_type')
plt.xlabel('Working Day')
plt.ylabel('Rental count')
plt.show()
```




- Demand from registered users is higher during working days
- Demand from casual users is higher during holidays

0.2.4 Weather

```
In [31]: plt.figure(figsize=(10,5))
sns.barplot(data=train,x='weather',y='count',ci=None)
plt.title('')
ax = plt.axes()
plt.xticks(rotation=0)
ax.set_xticklabels(['Clear','Mist','Light Snow/Rain','Heavy Rain/Thunderstorm'])
plt.title('Rentals by weather')
plt.xlabel('Weather Type')
plt.ylabel('Rental count')
plt.show()
```

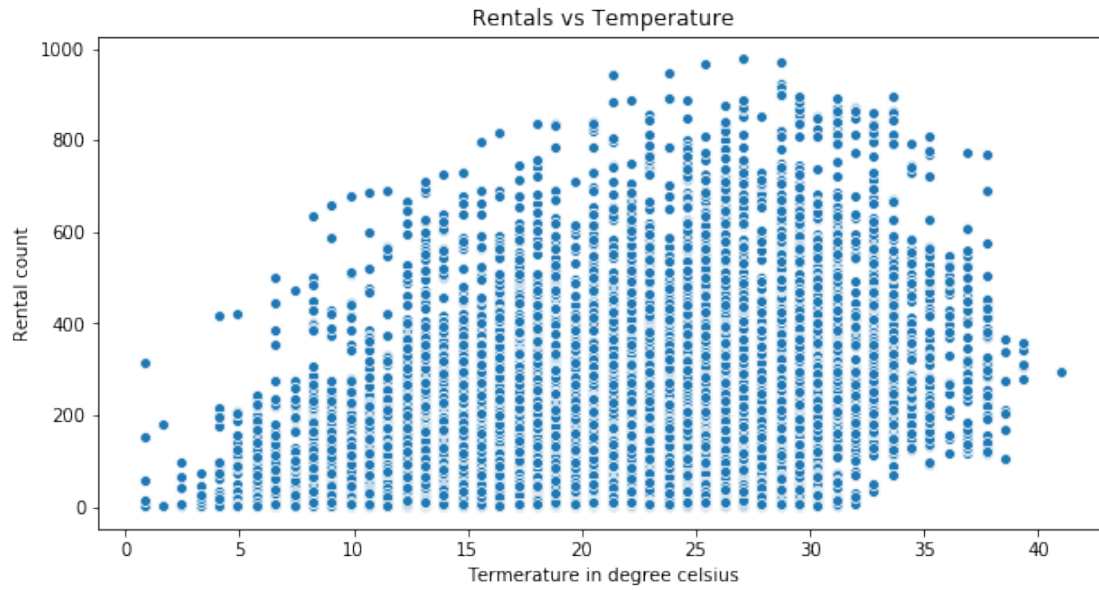
D:\Anaconda\lib\site-packages\matplotlib\cbook\deprecation.py:107: MatplotlibDeprecationWarning
 warnings.warn(message, mplDeprecation, stacklevel=1)




- People prefer to rent bikes during good weather.
- Interestingly, there is higher demand during Heavy Rains and Thunderstorms than during light rain or snow.

0.2.5 Temperature

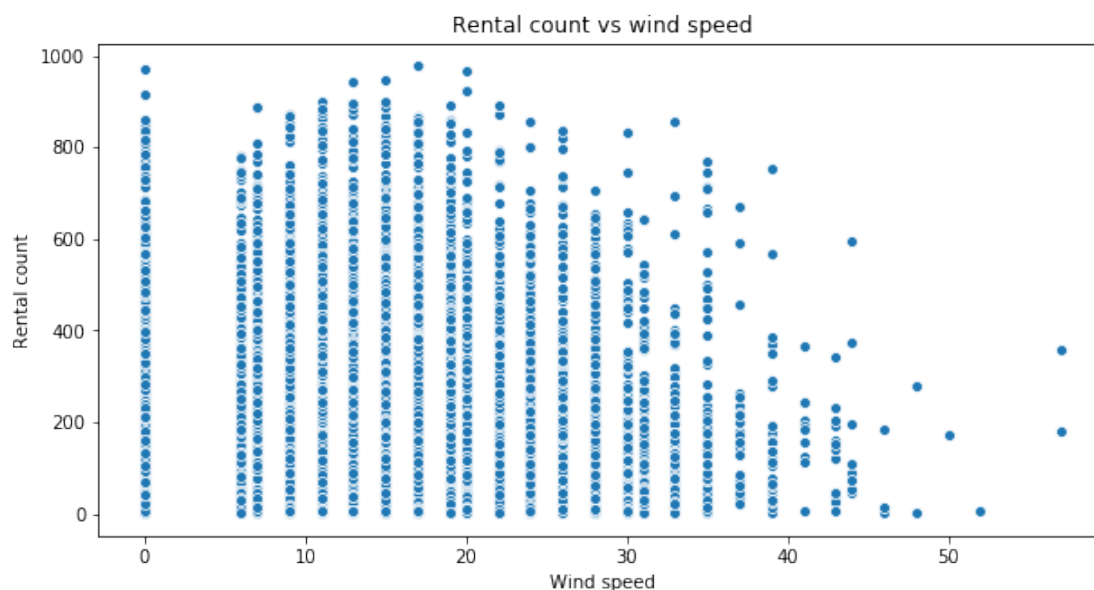
```
In [32]: plt.figure(figsize=(10,5))
sns.scatterplot(data=train,x='temp',y='count',ci=None)
plt.title('Rentals vs Temperature')
plt.xlabel('Temperature in degree celsius')
plt.ylabel('Rental count')
plt.show()
```



```
In [33]: plt.figure(figsize=(10,5))
sns.scatterplot(data=train,x='humidity',y='count',ci=None)
plt.title('Rental count vs Humidity')
plt.xlabel('Humidity')
plt.ylabel('Rental count')
plt.show()
```



```
In [34]: plt.figure(figsize=(10,5))
sns.scatterplot(data=train,x='windspeed',y='count',ci=None)
plt.title('Rental count vs wind speed')
plt.xlabel('Wind speed')
plt.ylabel('Rental count')
plt.show()
```



```
In [35]: train.drop(['weekday_name'],inplace=True,axis=1)
```

0.3 Feature Engineering

```
In [36]: generateColumnInfo(train)
```

```
Out[36]:
```

	ColumnName	NullCount	NonNullCount	NullPercent	UniqueValueCount	\
0	season	0	10886	0.0	4	
1	holiday	0	10886	0.0	2	
2	workingday	0	10886	0.0	2	
3	weather	0	10886	0.0	4	
4	temp	0	10886	0.0	49	
5	atemp	0	10886	0.0	60	
6	humidity	0	10886	0.0	89	
7	windspeed	0	10886	0.0	28	
8	casual	0	10886	0.0	309	
9	registered	0	10886	0.0	731	
10	count	0	10886	0.0	822	
11	date	0	10886	0.0	10886	
12	date_year	0	10886	0.0	2	
13	date_month	0	10886	0.0	12	

14	date_day	0	10886	0.0	19
15	date_weekday	0	10886	0.0	7
16	date_hour	0	10886	0.0	24

	DataType
0	int64
1	int64
2	int64
3	int64
4	float64
5	float64
6	int64
7	float64
8	int64
9	int64
10	int64
11	datetime64[ns]
12	int64
13	int64
14	int64
15	int64
16	int64

In [37]: generateColumnInfo(test)

Out [37]:	ColumnName	NullCount	NonNullCount	NullPercent	UniqueValueCount	\
0	datetime	0	6493	0.0	6493	
1	season	0	6493	0.0	4	
2	holiday	0	6493	0.0	2	
3	workingday	0	6493	0.0	2	
4	weather	0	6493	0.0	4	
5	temp	0	6493	0.0	49	
6	atemp	0	6493	0.0	65	
7	humidity	0	6493	0.0	79	
8	windspeed	0	6493	0.0	27	
9	date	0	6493	0.0	6493	
10	date_year	0	6493	0.0	2	
11	date_month	0	6493	0.0	12	
12	date_day	0	6493	0.0	12	
13	date_weekday	0	6493	0.0	7	
14	date_hour	0	6493	0.0	24	

	DataType
0	object
1	int64
2	int64
3	int64
4	int64

```


5         float64
6         float64
7         int64
8         float64
9    datetime64[ns]
10        int64
11        int64
12        int64
13        int64
14        int64

```

```

In [38]: plt.figure(figsize=(15,10))
         sns.heatmap(train.corr(),annot=True)
         plt.show()

```



- Date_hour shows the highest correlation with count.
- There is a very high co-relation between season and date_month and temp and atemp.
- I will drop season and atemp columns to avoid multicollinearity.

```

In [39]: train.drop(['season','atemp'],inplace=True,axis=1)
         test.drop(['season','atemp'],inplace=True,axis=1)

```


One hot encoding "weather"

```
In [40]: train = pd.get_dummies(train, columns=['weather'], prefix=['weather_'])
        test = pd.get_dummies(test, columns=['weather'], prefix=['weather_'])
```

The date column is no longer required.

```
In [41]: train.drop('date',inplace=True,axis=1)
        test.drop('date',inplace=True,axis=1)
```

```
In [42]: train['date_year'] = train['date_year'].astype("category")
        train['date_month'] = train['date_month'].astype("category")
        train['date_day'] = train['date_day'].astype("category")
        train['date_weekday'] = train['date_weekday'].astype("category")
        train['date_hour'] = train['date_hour'].astype("category")
```

0.4 Statistical Modeling

```
In [43]: X = train.drop(['count','casual','registered'],axis=1)
        y = train['count']
```

```
In [44]: X_train,X_valid, y_train,y_valid = train_test_split(X,y, test_size=0.3,random_state=1)
```

```
In [45]: res_alg = ['Random Forest','SVR','KNN','GBR','LR']
        res_rmsle = []
```

```
In [46]: from sklearn.ensemble import RandomForestRegressor
        from sklearn.svm import SVR
        from sklearn.ensemble import GradientBoostingRegressor
        from sklearn.neighbors import KNeighborsRegressor
        from sklearn.linear_model import LinearRegression
        from sklearn.model_selection import GridSearchCV
```

```
In [47]: #Random forest
        rfModel = RandomForestRegressor()
        rfModel.fit(X_train,np.log1p(y_train))
        rfPred = rfModel.predict(X_valid)
        np.sqrt(mean_squared_log_error(np.exp(rfPred),y_valid))
        res_rmsle.append(np.sqrt(mean_squared_log_error(np.exp(rfPred),y_valid)))
```

```
In [48]: #Support vector regression
        svrModel = SVR()
        svrModel.fit(X_train,np.log1p(y_train))
        svrPred = svrModel.predict(X_valid)
        res_rmsle.append(np.sqrt(mean_squared_log_error(np.exp(svrPred),y_valid)))
```

```
In [49]: #K-nearest neighbors
        knnModel = KNeighborsRegressor()
        knnModel.fit(X_train,np.log1p(y_train))
        knnPred = knnModel.predict(X_valid)
        res_rmsle.append(mean_squared_log_error(np.exp(knnPred),y_valid))
```

```
In [50]: #Gradient boosting regressor
gbModel = GradientBoostingRegressor()
gbModel.fit(X_train,np.log1p(y_train))
gbPred = gbModel.predict(X_valid)
gbPred = pd.DataFrame(gbPred)
gbPred = gbPred[0].apply(lambda x:0.0 if x<0 else x)
res_rmsle.append(np.sqrt(mean_squared_log_error(np.exp(gbPred),y_valid)))
```

```
In [51]: #Linear regression
lrModel = LinearRegression()
lrModel.fit(X_train,np.log1p(y_train))
lrPred = lrModel.predict(X_valid)
lrPred = pd.DataFrame(lrPred)
lrPred = lrPred[0].apply(lambda x:0.0 if x<0 else x)
res_rmsle.append(np.sqrt(mean_squared_log_error(np.exp(lrPred),y_valid)))
```

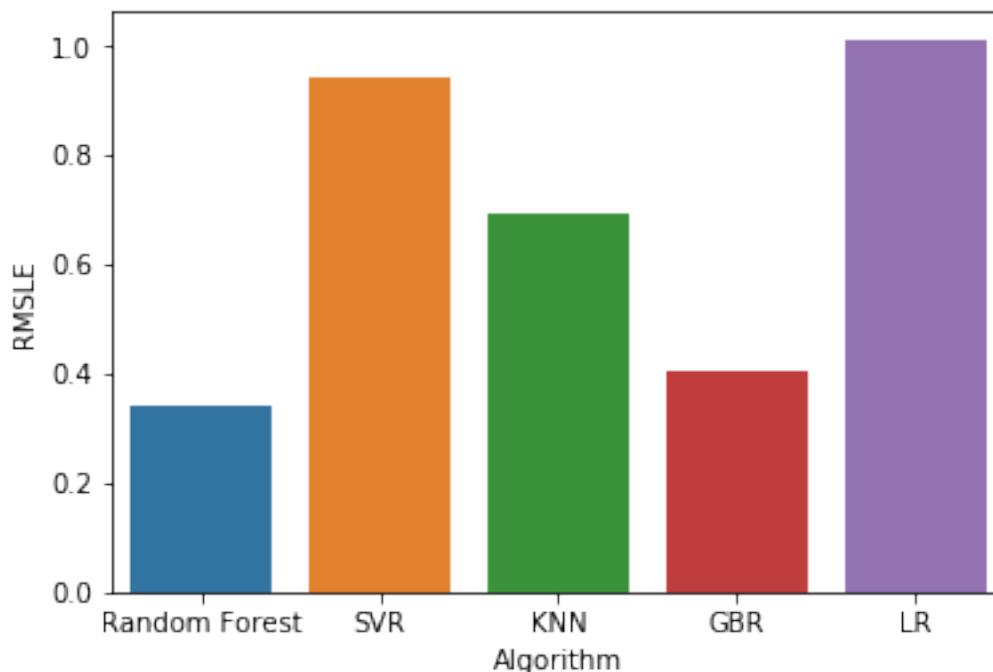
```
In [52]: results = pd.DataFrame({'Algorithm':res_alg,'RMSLE':res_rmsle})
results.sort_values('RMSLE')
```

```
Out [52]:
```

	Algorithm	RMSLE
0	Random Forest	0.338849
3	GBR	0.405260
2	KNN	0.691466
1	SVR	0.942840
4	LR	1.011953

```
In [53]: sns.barplot(data=results,x='Algorithm',y='RMSLE')
```

```
Out [53]: <matplotlib.axes._subplots.AxesSubplot at 0xfb2ec5e80>
```



Random forest gives the lowest RMSLE value. In the next section, I will tune the hyperparameters of the random forest model.

0.5 Tuning

0.5.1 Random Search

I will define a grid of hyperparameters and use Scikit-Learn's `RandomizedSearchCV` method to randomly sample from the grid, performing K-Fold CV with each combination of values. `RandomizedSearchCV` will not try every combination, but selects at random to sample a wide range of values. K-Fold CV reduces overfitting. I will try adjusting the following set of hyperparameters:

- `n_estimators` = number of trees in the forest
- `max_features` = max number of features considered for splitting a node
- `max_depth` = max number of levels in each decision tree
- `min_samples_split` = min number of data points placed in a node before the node is split
- `min_samples_leaf` = min number of data points allowed in a leaf node
- `bootstrap` = method for sampling data points (with or without replacement)

```
In [54]: n_estimators = [int(x) for x in np.linspace(start = 200, stop = 2000, num = 10)]
         max_features = ['auto', 'sqrt']
         max_depth = [int(x) for x in np.linspace(10, 110, num = 11)]
         max_depth.append(None)
         min_samples_split = [2, 5, 10]
         min_samples_leaf = [1, 2, 4]
         bootstrap = [True, False]
```

```
In [55]: random_grid = {'n_estimators': n_estimators,
                        'max_features': max_features,
                        'max_depth': max_depth,
                        'min_samples_split': min_samples_split,
                        'min_samples_leaf': min_samples_leaf,
                        'bootstrap': bootstrap}
```

```
In [56]: from sklearn.model_selection import RandomizedSearchCV
```

```
In [57]: rf = RandomForestRegressor()
         rf_random = RandomizedSearchCV(estimator = rf, param_distributions = random_grid, n_i
```

```
In [58]: rf_random.fit(X_train,np.log1p(y_train))
```

Fitting 3 folds for each of 100 candidates, totalling 300 fits

```
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed: 1.3min
[Parallel(n_jobs=-1)]: Done 184 tasks     | elapsed: 7.1min
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed: 11.4min finished
```

```
Out [58]: RandomizedSearchCV(cv=3, error_score='raise',
                             estimator=RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=
                             max_features='auto', max_leaf_nodes=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=1,
                             oob_score=False, random_state=None, verbose=0, warm_start=False),
                             fit_params=None, iid=True, n_iter=100, n_jobs=-1,
                             param_distributions={'n_estimators': [200, 400, 600, 800, 1000, 1200, 1400,
                             pre_dispatch='2*n_jobs', random_state=42, refit=True,
                             return_train_score='warn', scoring='neg_mean_squared_log_error',
                             verbose=1)
```

```
In [59]: rf_random.best_params_
```

```
Out [59]: {'bootstrap': True,
           'max_depth': 100,
           'max_features': 'auto',
           'min_samples_leaf': 1,
           'min_samples_split': 2,
           'n_estimators': 1400}
```

0.5.2 Grid Search CV

Now that we know the range for each hyperparameter, I will fine tune the hyperparameters with GridSearchCV. GridSearchCV will perform KFold Cross validation on each combination.

```
In [60]: param_grid = {
           'bootstrap': [True],
           'max_depth': [100,120,150],
           'max_features': ['auto'],
           'min_samples_leaf': [1,2],
           'min_samples_split': [2,3,4],
           'n_estimators': [500,800,1000 ]
         }
```

```
In [61]: from sklearn.model_selection import GridSearchCV
```

```
In [62]: rf = RandomForestRegressor()
           rf_grid = GridSearchCV(estimator =rf,param_grid=param_grid,scoring='neg_mean_squared_
```

```
In [63]: rf_grid.fit(X_train,np.log1p(y_train))
           pred=rf_grid.predict(X_valid)
```

Fitting 3 folds for each of 54 candidates, totalling 162 fits

```
[Parallel(n_jobs=1)]: Done 162 out of 162 | elapsed: 29.6min finished
```

```
In [64]: print((np.sqrt(mean_squared_log_error(np.exp(pred),y_valid))))
```

0.317396652773138

Tuning has led to a decrease in the RMSLE value. Let us look at the final hyperparameters:


```
In [65]: rf_grid.best_params_
```

```
Out[65]: {'bootstrap': True,
          'max_depth': 150,
          'max_features': 'auto',
          'min_samples_leaf': 1,
          'min_samples_split': 3,
          'n_estimators': 800}
```

```
In [67]: f_imp = pd.DataFrame({'Feature':X_train.columns,'Importance':rf_grid.best_estimator_.feature_importances_})
```

```
In [68]: plt.figure(figsize=(10,6))
          sns.barplot(data=f_imp.sort_values(by='Importance',ascending=False),x='Importance',y=
          plt.title('Feature Importance')
```

```
Out[68]: Text(0.5,1,'Feature Importance')
```



0.6 Generate predictions on Test

```
In [69]: testPred=rf_grid.predict(test.drop('datetime',axis=1))
         testPred= np.exp(testPred)
         d={'datetime':test['datetime'],'count':testPred}
         ans=pd.DataFrame(d)
         ans.to_csv('answer.csv',index=False)
```